

# FastAPI Message Broker

A Minimal In-Memory Message Broker System

Similar to Apache Kafka

Built with Python & FastAPI

# What We Are Building

## System Components & Architecture

---

### Three Core Components

- ✓ **Broker** - FastAPI server that stores messages per topic
- ✓ **Producer** - Python script that publishes messages via HTTP
- ✓ **Consumer** - Python script that polls and reads messages

### Key Constraints

- Data exists only in RAM (in-memory)
- One message list per topic
- No partitions or replication
- Each consumer has independent offset tracking

# Project Structure

Organised for Clarity & Maintainability

```
mini-kafka/ ├── broker/ | └─ main.py # FastAPI app (the broker) ├── producer/ | └─ main.py # Producer script ├── consumer/ | └─  
main.py # Consumer script ├── shared/ | └─ schemas.py # Pydantic models shared by all apps ├── requirements.txt └─ README.md
```

## Why This Split?

- **broker/** - Web server with REST endpoints
- **producer/consumer/** - Plain Python HTTP clients
- **shared/** - Avoids duplicating request/response models

# Data Models (Pydantic)

Type-Safe Request & Response Schemas

Model	Purpose	Fields
<b>TopicRegistration</b>	Ensure topic exists	{ topic: str }
<b>ProducerRegistration</b>	Register a producer	{ topic: str, producer_id?: str }
<b>ConsumerRegistration</b>	Register a consumer	{ topic: str, consumer_id?: str }
<b>PublishRequest</b>	Publish a message	{ topic: str, value: str, key?: str }
<b>ConsumeRequest</b>	Request next message	{ consumer_id: str }
<b>Message</b>	Message response	{ topic: str, offset: int, value: str, key?: str }

# API Endpoints

RESTful Interface with Auto-Generated Swagger Docs

Method	Endpoint	Purpose
GET	/	Welcome message
GET	/health	Health check
POST	/topics/register	Create topic array if missing
POST	/producers/register	Acknowledge producer (logging/demo)
POST	/consumers/register	Create consumer, set offset to 0
POST	/produce	Append message, assign offset
POST	/consume	Return next message, increment offset
GET	/stats	Return counts per topic and consumer offsets

# Consumer Workflow

## Polling & Reading Messages

```
Consumer                               Broker
| POST /topics/register |             |
|----->| ensure topic exists
|
| POST /consumers/register (topic)
|----->| {consumer_id, offset=0}
|
Repeat (infinite loop):
| POST /consume {consumer_id}
|----->|
| <-----| 200 {message} or
|         | 204 (no content)
```

## Response Handling

- **200 OK** - Message available, print and continue
- **204 No Content** - No new messages, sleep for poll interval

# Understanding Offsets

## Independent Tracking Per Consumer

### Example: Topic "demo" with messages [0, 1, 2]

- Consumer **alice** starts with offset = 0
- First /consume → returns message at index 0, offset becomes 1
- Next /consume → returns index 1, offset becomes 2
- If no message at index 2, broker returns 204

### Multiple Consumers

Each consumer's offset is **independent**. A second consumer **bob** also starts at offset 0 and will receive all messages from the beginning.

**Key Point:** Offsets are per-consumer, enabling parallel independent consumption

# Key Technical Features

What Makes This System Robust

## Architecture Benefits

- ✓ Async operations with FastAPI
- ✓ Thread-safe with per-topic locks
- ✓ Type-safe with Pydantic models
- ✓ Auto-generated API documentation
- ✓ Modular component separation

## Message Guarantees

- ✓ Message ordering within topics
- ✓ Independent consumer offsets
- ✓ Configurable polling intervals
- ✓ Built-in monitoring via /stats
- ✓ HTTP-based, language-agnostic



## Common Pitfalls & Solutions

### Troubleshooting Guide

Issue	Cause	Solution
204 No Content	No new messages available	Not an error - keep polling
404 Topic Not Found	Topic not registered	Call /topics/register first
Data Lost on Restart	In-memory storage by design	Add disk persistence if needed
Consumer Conflicts	Duplicate consumer IDs	Use unique --consumer-id per consumer
Low Throughput	Single-process limitations	Add partitions, batching, or workers

# How to Extend the System

## Future Enhancement Opportunities

---

### Performance & Scale

- **Partitions** - Multiple shards per topic
- **Batching** - Produce/consume multiple messages
- **Backpressure** - Max topic size limits
- **Long-polling** - Reduce polling overhead

### Enterprise Features

- **Durability** - SQLite or file persistence
- **Consumer Groups** - Load balancing
- **Auth & ACLs** - Security controls
- **Offset Commits** - At-least-once semantics

# Summary

## What We've Built

---

### Core Achievement

A **lightweight, maintainable message broker** system with clear separation of concerns, built on modern Python async principles.

### Key Takeaways

- ✓ Simple yet functional message broker architecture
- ✓ Async FastAPI for high-performance HTTP endpoints
- ✓ Independent offset tracking for multiple consumers
- ✓ Type-safe with Pydantic, developer-friendly with Swagger
- ✓ Foundation for learning distributed systems concepts

---

**Perfect for:** Learning, prototyping, and understanding message broker fundamentals

# Thank You

Questions?

FastAPI Message Broker  
Built with Python, FastAPI & Pydantic