

# FastAPI Message Broker

A Minimal In-Memory Message Broker System

Similar to Apache Kafka

Built with Python & FastAPI

# What We Are Building

## System Components & Architecture

---

### Three Core Components

- ✓ **Broker** - FastAPI server that stores messages per topic
- ✓ **Producer** - Python script that publishes messages via HTTP
- ✓ **Consumer** - Python script that polls and reads messages

### Key Constraints

- Data exists only in RAM (in-memory)
- One message list per topic
- No partitions or replication
- Each consumer has independent offset tracking

# Project Structure

Organised for Clarity & Maintainability

```
mini-kafka/ ├── broker/ | └─ main.py # FastAPI app (the broker) ├── producer/ | └─ main.py # Producer script ├── consumer/ | └─  
main.py # Consumer script ├── shared/ | └─ schemas.py # Pydantic models shared by all apps ├── requirements.txt └─ README.md
```

## Why This Split?

- **broker/** - Web server with REST endpoints
- **producer/consumer/** - Plain Python HTTP clients
- **shared/** - Avoids duplicating request/response models

# Data Models (Pydantic)

Type-Safe Request & Response Schemas

| Model                       | Purpose              | Fields   |
|-----------------------------|----------------------|--|
| <b>TopicRegistration</b>    | Ensure topic exists  | { topic: str }                                     |
| <b>ProducerRegistration</b> | Register a producer  | { topic: str, producer_id?: str }                  |
| <b>ConsumerRegistration</b> | Register a consumer  | { topic: str, consumer_id?: str }                  |
| <b>PublishRequest</b>       | Publish a message    | { topic: str, value: str, key?: str }              |
| <b>ConsumeRequest</b>       | Request next message | { consumer_id: str }                               |
| <b>Message</b>              | Message response     | { topic: str, offset: int, value: str, key?: str } |

# API Endpoints

RESTful Interface with Auto-Generated Swagger Docs

| Method | Endpoint            | Purpose                                      |
|--------|---------------------|--|
| GET    | /                   | Welcome message                              |
| GET    | /health             | Health check                                 |
| POST   | /topics/register    | Create topic array if missing                |
| POST   | /producers/register | Acknowledge producer (logging/demo)          |
| POST   | /consumers/register | Create consumer, set offset to 0             |
| POST   | /produce            | Append message, assign offset                |
| POST   | /consume            | Return next message, increment offset        |
| GET    | /stats              | Return counts per topic and consumer offsets |

# Consumer Workflow

## Polling & Reading Messages

```
Consumer                               Broker
| POST /topics/register |               |
|----->| ensure topic exists
|
| POST /consumers/register (topic)
|----->| {consumer_id, offset=0}
|
Repeat (infinite loop):
| POST /consume {consumer_id}
|----->|
| <-----| 200 {message} or
|         | 204 (no content)
```

## Response Handling

- **200 OK** - Message available, print and continue
- **204 No Content** - No new messages, sleep for poll interval

# Understanding Offsets

## Independent Tracking Per Consumer

### Example: Topic "demo" with messages [0, 1, 2]

- Consumer **alice** starts with offset = 0
- First /consume → returns message at index 0, offset becomes 1
- Next /consume → returns index 1, offset becomes 2
- If no message at index 2, broker returns 204

### Multiple Consumers

Each consumer's offset is **independent**. A second consumer **bob** also starts at offset 0 and will receive all messages from the beginning.

**Key Point:** Offsets are per-consumer, enabling parallel independent consumption

# Key Technical Features

What Makes This System Robust

---

## Architecture Benefits

- ✓ Async operations with FastAPI
- ✓ Thread-safe with per-topic locks
- ✓ Type-safe with Pydantic models
- ✓ Auto-generated API documentation
- ✓ Modular component separation

## Message Guarantees

- ✓ Message ordering within topics
- ✓ Independent consumer offsets
- ✓ Configurable polling intervals
- ✓ Built-in monitoring via /stats
- ✓ HTTP-based, language-agnostic



## Common Pitfalls & Solutions

### Troubleshooting Guide

| Issue                | Cause                       | Solution                              |
|----------------------|-----------------------------|---------------------------------------|
| 204 No Content       | No new messages available   | Not an error - keep polling           |
| 404 Topic Not Found  | Topic not registered        | Call /topics/register first           |
| Data Lost on Restart | In-memory storage by design | Add disk persistence if needed        |
| Consumer Conflicts   | Duplicate consumer IDs      | Use unique --consumer-id per consumer |
| Low Throughput       | Single-process limitations  | Add partitions, batching, or workers  |

# How to Extend the System

## Future Enhancement Opportunities

---

### Performance & Scale

- **Partitions** - Multiple shards per topic
- **Batching** - Produce/consume multiple messages
- **Backpressure** - Max topic size limits
- **Long-polling** - Reduce polling overhead

### Enterprise Features

- **Durability** - SQLite or file persistence
- **Consumer Groups** - Load balancing
- **Auth & ACLs** - Security controls
- **Offset Commits** - At-least-once semantics

# Backpressure, Limits & Safety

## *Production-Ready Message Broker Protection*

### 1 Rate Limiting

Implement **token bucket algorithm** per IP/topic to protect broker from overwhelming traffic. Prevents resource exhaustion and ensures fair access for all clients.

### 2 Payload Size Validation

Enforce **maximum payload size** checks on `PublishRequest` to prevent memory issues and ensure consistent performance across the system.

### 3 Queue Management

Set **queue limits per topic** with configurable strategies: drop oldest messages or reject new ones with `429 Too Many Requests` response.

# Backpressure, Limits & Safety

*Production-Ready Message Broker Protection*

## 4 Idempotency Keys

Accept `message_id` from producers and maintain a **recent ID cache** for deduplication. Prevents duplicate message processing and ensures exactly-once semantics.

## 5 Dead-Letter Queue (DLQ)

Automatically move messages to `topic.DLQ` when consumers fail to commit after **N attempts or timeout threshold**. Enables manual inspection and prevents message loss.

## ✓ Implementation Priority

These safety mechanisms are **critical for production readiness** and should be implemented before scaling to handle high-volume traffic.

# Summary

## What We've Built

---

### Core Achievement

A **lightweight, maintainable message broker** system with clear separation of concerns, built on modern Python async principles.

### Key Takeaways

- ✓ Simple yet functional message broker architecture
- ✓ Async FastAPI for high-performance HTTP endpoints
- ✓ Independent offset tracking for multiple consumers
- ✓ Type-safe with Pydantic, developer-friendly with Swagger
- ✓ Foundation for learning distributed systems concepts

---

**Perfect for:** Learning, prototyping, and understanding message broker fundamentals

# Thank You

Questions?

FastAPI Message Broker  
Built with Python, FastAPI & Pydantic